

Aditi Reddy Doma

domaaditi@gmail.com | 917-337-7045 | [linkedin.com/in/aditireddydoma](https://www.linkedin.com/in/aditireddydoma) | [GitHub](#) | [Portfolio](#)

EDUCATION

Masters of Science in Data Science

Rochester Institute of Technology (GPA: 4.0/4.0)

Expected: December 2026

Rochester, New York

Relevant Coursework: Applied Statistics(A), Database Management(A), Visual Analytics(A), Software Construction(A)

Bachelor of Technology in Artificial Intelligence

Mahindra University (GPA: 3.56/4.0)

July 2020 - May 2024

Hyderabad, India

Relevant Coursework: Machine Learning(A-), Data Structures & Algorithms(A), Generative Artificial Intelligence(A-), Operating Systems(A-)

SKILLS

- **Programming Languages:** Python, Java, R, SQL
- **Cloud Platforms:** AWS, Azure, Snowflake, Databricks, Kafka, Microsoft Fabric
- **Data Engineering:** ETL, CDC, Data Pipelines, Apache Airflow, PySpark, SQL/NoSQL Databases
- **BI & Analytics:** Tableau, Power BI, Looker, Google Analytics, Streamlit
- **ML & AI:** TensorFlow, PyTorch, Scikit-Learn, BERT, LLMs, RAG, HuggingFace, CLIP, FAISS, YOLOv8, MediaPipe, OpenCV, Neo4j
- **Tools & Technologies:** Git, KQL, Docker, Jupyter

WORK EXPERIENCE

Artificial Intelligence Engineer Intern

Johnson Controls

August 2026 - Present

Sunnyvale, California

- **Developing the Control Spec Assistant**, an AI agent using **hybrid RAG (BGE-large + BM25, Qdrant) and GPT-5.1** to deliver **instant, cited spec answers**, targeting **157 hrs/week** that **18 controls engineers** spend manually searching documents.
- Architecting **Agent Q**, an **agentic AI chatbot** that **converts engineers' natural-language questions into SQL** over the Snowflake chiller sensor database, **delivering instant performance insights without writing code, accelerated with Claude Code**.

Data Engineering Intern

Johnson Controls

June 2026 - August 2026

York, Pennsylvania

- Built an **ETL pipeline** from **Microsoft Fabric to Snowflake** using **KQL and Python**, centralizing **17M+ chiller sensor records**.
- Defined **15+ equipment health KPIs** and built **Tableau dashboards** for **cross-functional teams** to monitor equipment reliability.
- **Won the Johnson Controls Summer Innovation Challenge** among **150+ North American interns** with an **AI-powered Predictive Sales Optimization Engine** recommending bid strategies from historical win/loss data.

Data Science Research Assistant

iCitizen

May 2025 - Present

Rochester, New York

- **Web scraped and analyzed 500K+ unstructured congressional profiles** from Congress.gov and Bioguide websites using APIs, **transforming unstructured data into structured datasets** stored in **AWS DynamoDB** for **real-time updates**.
- Designed a dynamic U.S. map using **D3.js**, visualizing senators, their sponsored bills and co-sponsors by state.
- **Increased user engagement by 27.2%** on the iCitizen app by building a **GraphRAG chatbot on Neo4j**, reducing hallucinations.

Data Analyst

StandardWings Technologies Pvt. Ltd.

August 2023 - August 2024

Nashik, India

- Extracted, cleaned and integrated **multi-source data** (GPS, sensor, Salesforce CRM, SAP, Excel) from **120+ vehicles** into **AWS S3** using **SQL and Python**, enabling **Power BI** fleet analysis that identified **37% fuel cost reduction opportunities**.
- Collaborated with **cross-functional teams** throughout the project lifecycle to define **21 key KPIs** aligned with business objectives.

Machine Learning Intern

Cloud4C

June 2023 - July 2023

Hyderabad, India

- Achieved **over 90% accuracy** in a **sentiment analysis model** using **BERT** to analyze customer feedback.
- **Performed large-scale text analysis** on customer feedback using **Databricks** with **PySpark** for data processing.
- Identified **top 5 recurring customer concerns** using **unsupervised learning (K-Means clustering)** on BERT embeddings.

PROJECTS

e-CommMart: e-Commerce Data Platform | AWS, MariaDB, MongoDB, Python, ETL, CDC, NoSQL, Vector Search

[GitHub](#)

- Engineered a hybrid data platform using **Python/SQL ETL**, **MariaDB**, and **MongoDB** on **AWS EC2**, implementing **CDC-based synchronization** between relational and document stores for downstream analytics.
- Deployed a highly available **3-node MariaDB Galera Cluster** with replication, fault tolerance, RBAC, data-quality logging, and secure access controls.

Melody: Spotify Analytics Platform | HuggingFace, Snowflake, Kafka, Airflow, PySpark, Tableau, Streamlit, Generative AI

[GitHub](#)

- Implemented **ETL pipeline** using **Apache Airflow**, **Kafka** and **PySpark** to stream Spotify data into **Snowflake** for analysis.
- Designed Streamlit dashboard with **Phi-3.5-Instruct** for playlist recommendations and Tableau dashboard for music analytics.